



# ML 2023 Project Slides

Authors:

Chiara Checchetti (583054) - [c.cecchetti@studenti.unipi.it](mailto:c.cecchetti@studenti.unipi.it)

Nicola Emmolo (672926) - [n.emmolo@studenti.unipi.it](mailto:n.emmolo@studenti.unipi.it)

Team name: gradient\_decent

Master degree: MSc in Computer Science - Artificial Intelligence

Date: 31/01/2024

Type of project: A

# OBJECTIVES

- Implementation of an MLP capable of performing binary classification and linear regression tasks
- Optimization of the various losses considered and metrics used to evaluate the ability to generalize and predict on new data (accuracy and MEE respectively).
- Analyses of the best models obtained.
- Solve the MONK problems and obtain the output of the blind test for the CUP

# YOUR CONTRIBUTIONS

- Libraries: Numpy / Pandas / Matplotlib / Itertools
- Implementation
  - Multi Layer Perceptron with one hidden layer
  - Output: 1 unit for MONK, 3 units for CUP
  - Activation Functions
    - Monk: Sigmoid (hidden), TanH (output)
    - CUP: ReLU / Leaky ReLU / TanH (hidden), Identity (output)
  - Training: Minibatch Stochastic Gradient Descent and Backpropagation
  - Weights Initialization: Random Uniform
  - Bias Initialization: All 1
  - Regularization: Tikhonov Regularization (L2) + Early Stopping (only for CUP)
- Preprocessing
  - Monk: One Hot Encoding and preliminary statistical tests (to check for a normal distribution)
- Validation
  - Hold Out & K Fold Cross Validation ( $K = 5$ )

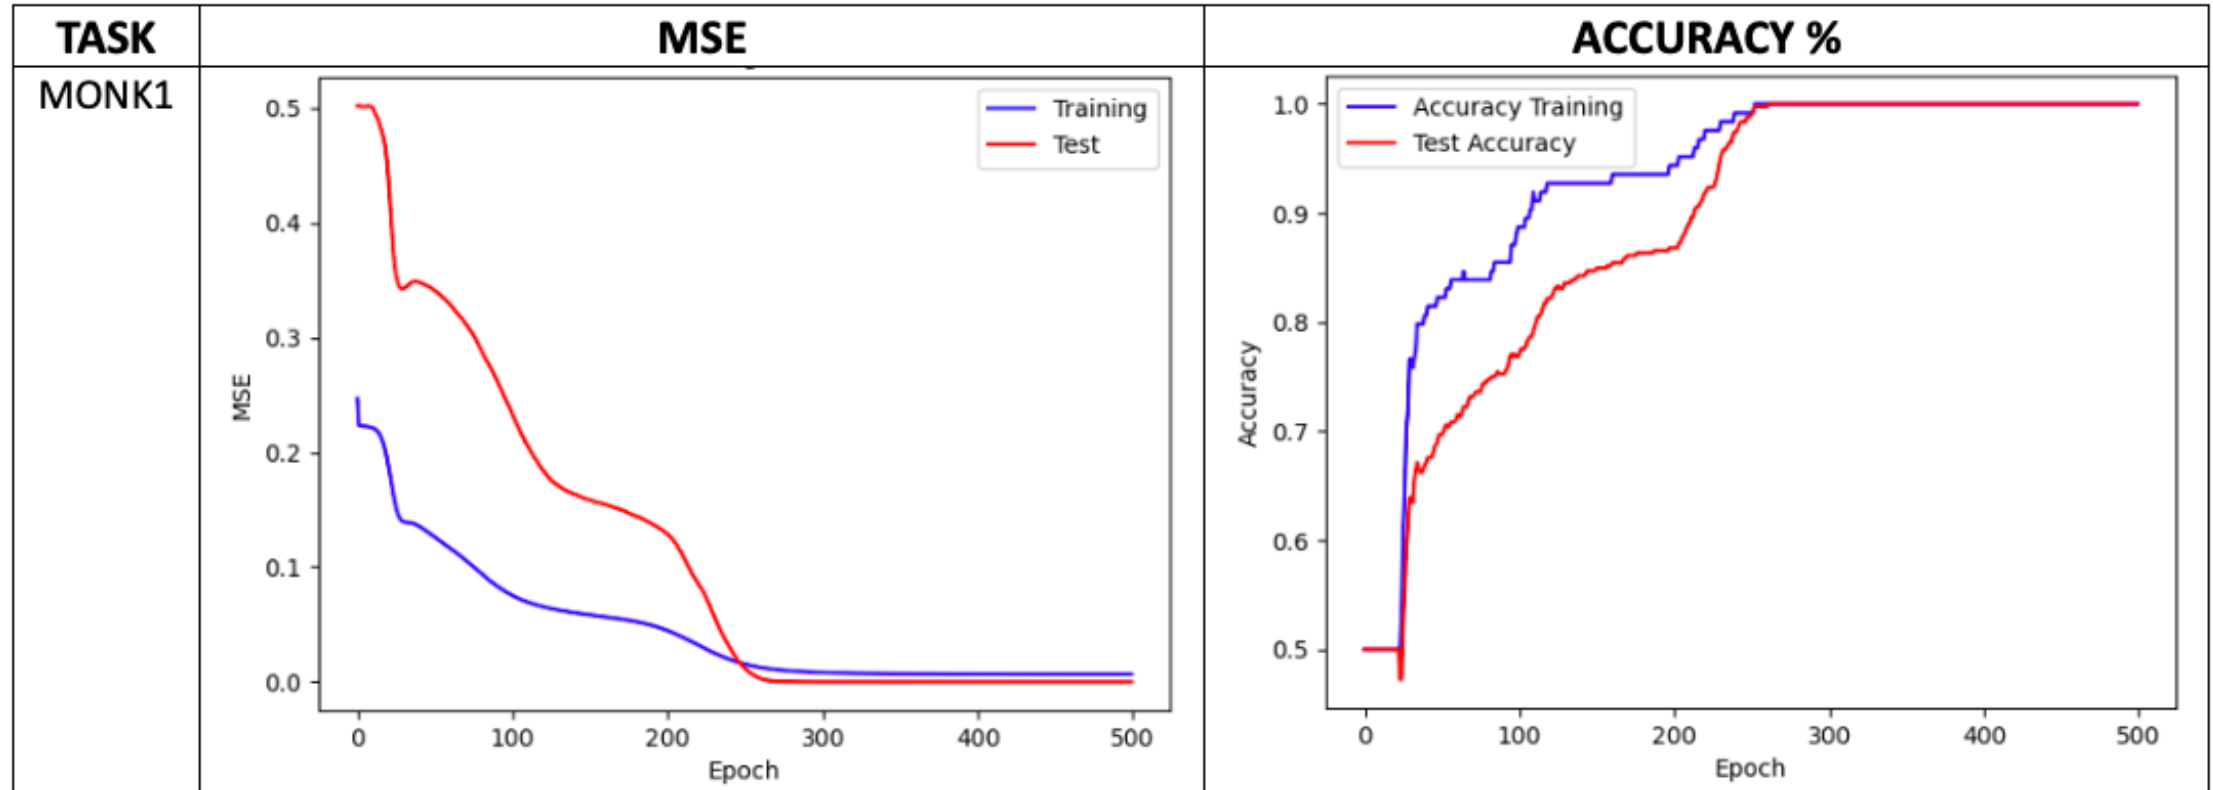
# MONK RESULTS

- To understand best initializations:
  - Small Grid Search, Theory, Personalized trials
- Fixed for every MONK:
  - **Epochs: 500 / Batch Size: 1 / Momentum: 0.9 / Weight Initialization Limits: [-0.2, 0.2]**


TASK	Hidden Units	Learning Rate	Lambda	MSE (TR/TS)	Accuracy (TR/TS)
<b>MONK1</b>	3	0.05	0	0.006632 / 0.0	100% / 100%
<b>MONK2</b>	4	0.4	0	0.000155 / 0.0	100% / 100%
<b>MONK3</b>	2	0.55	0	0.01626 / 0.06779	99.18% / 93.51%
<b>MONK3 (reg.)</b>	3	0.1	0.001	0.03467 / 0.02419	95.90% / 97.91%

Mean MSE and Accuracy over 5 trials

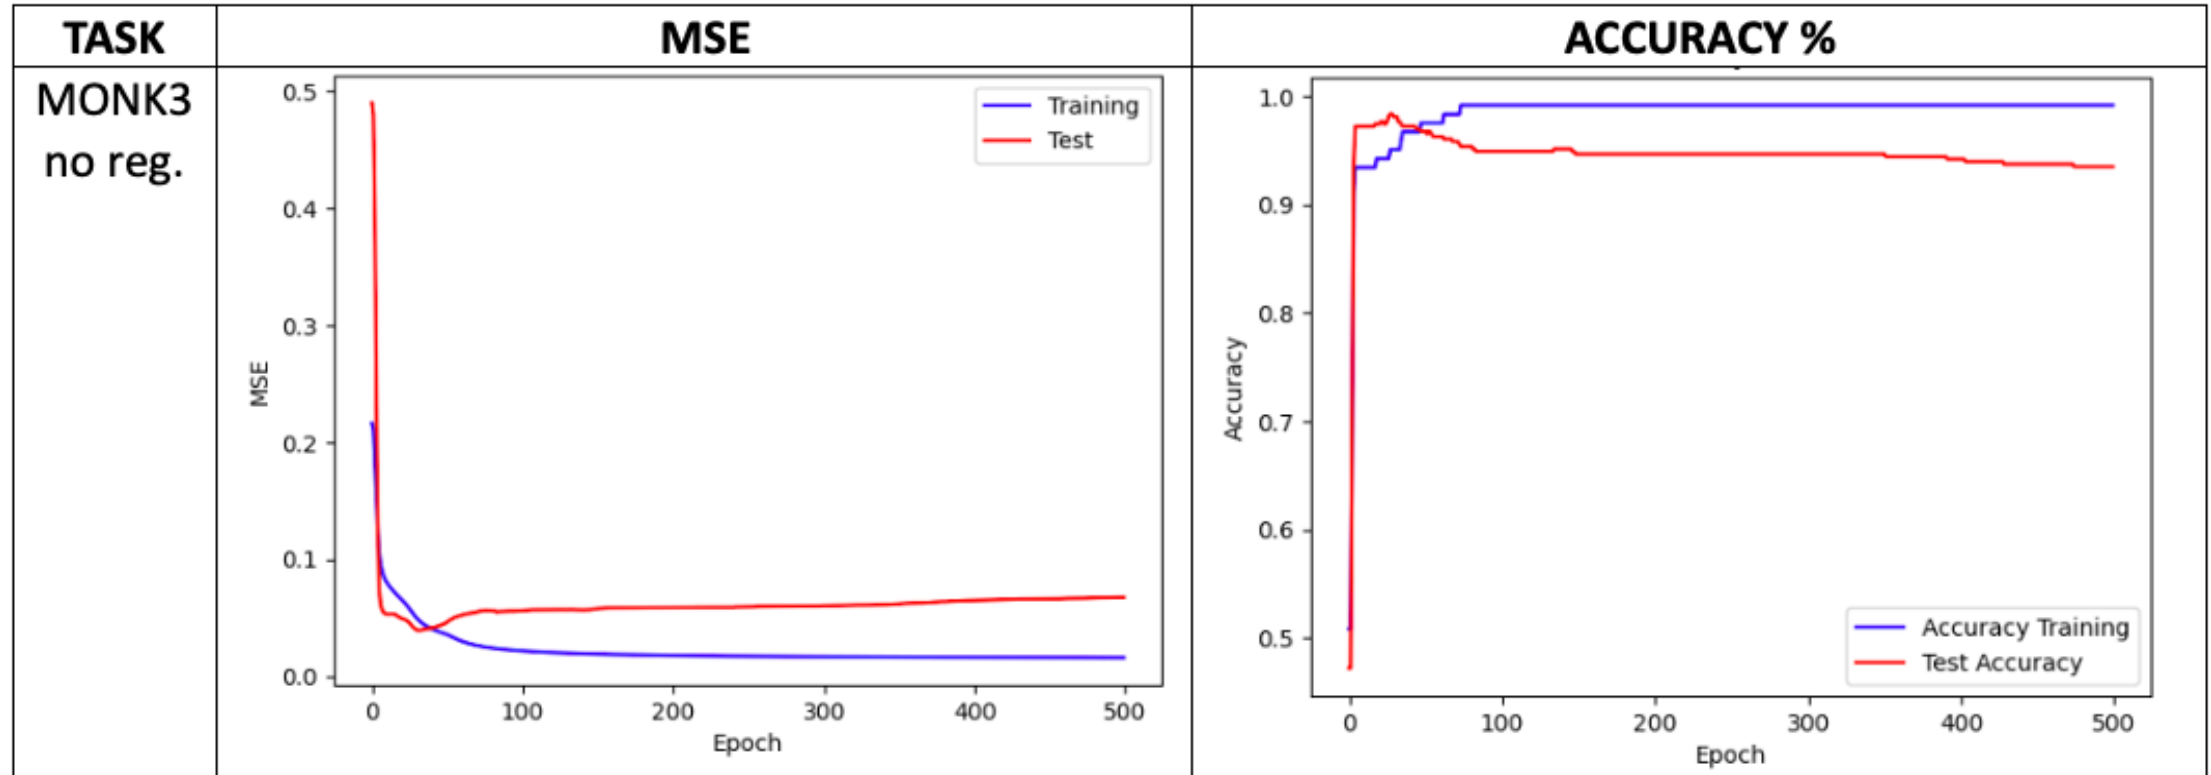
# MONK RESULTS




# MONK RESULTS



# MONK RESULTS



# MONK RESULTS





# MONK RESULTS: LIST OF THE HYPERPARAMETERS

- ***hidden\_size***: number of units of the hidden layer
- ***learning\_rate***: learning rate (eta)
- ***epochs***: the number of epochs
- ***batch\_size***: the mini-batch size (if 1 then online, else batch)
- ***momentum***: the momentum coefficient (alpha)
- ***lambda\_reg***: the regularization coefficient (for L2)
- ***w\_init\_limit***: the range for the initial weights

# CUP VALIDATION SCHEMA: DATA SPLITTING

- 80% Training+Validation=Development Set (800 examples)
  - K-Fold (K=5)
    - 4 Fold for Training (640 examples), 1 Fold for Validation (160 examples)
  - Hold-Out
    - 80% Training (640 examples), 20% Validation (160 examples)
- 20% Internal Test (200 examples)
- Retraining on Development Set

# CUP VALIDATION SCHEMA: MODEL SELECTION

- Screening Phase
  - High number of hidden units (universal theorem)
  - Initially using a coarse grid search (Hold-Out Validation), then with a finer grid search based on the results obtained with the first one

## COARSE

HYPERPARAMETER TYPE	RANGE
<i>hidden_size</i>	(40, 50) step 1
<i>learning_rate</i>	(0.05, 0.15) step 0.01
<i>epochs</i>	(400, 800) step 100
<i>batch_size</i>	[200, 400, 600, len_data]
<i>momentum</i>	(0.1, 0.2) step 0.01
<i>lambda_reg</i>	[0.000001, 0.00001, 0.0001]
<i>w_init_limit</i>	[[-0.5, 0.5], [-0.6, 0.6], [-0.7, 0.7]]

## FINE

HYPERPARAMETER TYPE	RANGE
<i>hidden_size</i>	[47]
<i>learning_rate</i>	(0.01, 0.1) step 0.001
<i>epochs</i>	(500, 600) step 100
<i>batch_size</i>	[200]
<i>momentum</i>	(0.1, 0.2) step 0.01
<i>lambda_reg</i>	[0.0001]
<i>w_init_limit</i>	[[-0.7, 0.7]]

# CUP VALIDATION SCHEMA: CUP RESULTS


- Coarse Grid Search with Hold Out
  - 10 hours using MSI 12th Gen Intel(R) Core(TM) i7-12650H 2.30 GHz
  - Number of combinations: 239580
  - Initial activation function: ReLU (hidden), Identity (output)
  - Patience for early stopping: 20
- With 5 best model we understand the range for the Fine Grid Search (combinations: 2000), and we:
  - Compare Hold-Out and K-fold
    - Similar results but K-Fold is better with smaller dataset
  - Compare ReLU / LeakyReLU / TanH on hidden layer
    - ReLU and Leaky ReLU seems better than TanH

# CUP RESULTS

- Comparison of different combinations of hyperparameters to check the differences and similarities of the results (plots and hyperparameters used are in the appendix)
- Final Model with K-Fold and ReLU (hidden)
  - Select best 10 model on Fine Grid Search, basing on validation error
  - Compare their plot
  - Compare their test error

# CUP RESULTS

- *hidden\_size*: 47
- *learning\_rate*: 0.08
- *epochs*: 600
- *batch\_size*: 200
- *momentum*: 0.11
- *lambda\_reg*: 1e-05
- *w\_init\_limit*: [-0.7, 0.7]
  
- Development MEE
  - 1.6379876671343174
- Internal Test MEE
  - 1.8067147539347013
- Validation MEE
  - 1.788410259793571



# DISCUSSION

- What we found significant / What we learned
  - Understand better the theoretical aspect creating the network from scratch
  - Understand the importance of validation and test part
  - How a small change on hyperparameter results in different solutions
  - How to avoid problems with small datasets
    - MONK: Learning rate small and momentum high
    - MONK: Online better than Batch
  - CUP: Using Regularization, still exploit Early Stopping (due to high chances of overfitting)

# CONCLUSION

- Compute the predictions for blind test, and save on file "gradient\_decent\_ML-CUP23-TS.csv"
- According to our Internal Test, the expected loss on Blind Test is around 1.80671475 (MEE)




# BIBLIOGRAPHY

- A. Micheli: "Derivation of the Back-propagation learning based algorithm"
- Monk Dataset
  - <https://archive.ics.uci.edu/dataset/70/monk+s+problems>
- A. Geron: "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow 3e: Concepts, Tools, and Techniques to Build Intelligent Systems"


# APPENDIX

<b>MODEL</b>	<b><i>hidden_size</i></b>	<b><i>learning_rate</i></b>	<b><i>epochs</i></b>	<b><i>batch_size</i></b>	<b><i>momentum</i></b>	<b><i>lambda_reg</i></b>	<b><i>w_init_limit</i></b>
1	100	0.1	600	200	0.1	0.0001	[-0.7, 0.7]
2	4	0.1	600	200	0.1	0.0001	[-0.7, 0.7]
3	30	0.9	600	200	0.1	0.0001	[-0.7, 0.7]
4	30	0.1	600	1	0.1	0.0001	[-0.7, 0.7]
5	30	0.1	600	200	0.9	0.0001	[-0.7, 0.7]
6	30	0.2	600	200	0.1	0.0001	[-0.1, 0.1]
7	30	0.1	600	16	0.1	0.001	[-0.5, 0.5]
8	30	0.1	600	4	0.1	0.0001	[-0.5, 0.5]
9	30	0.2	600	200	0.1	0.001	[-0.20, 0.20]
10	10	0.05	500	200	0.5	0.0001	[-0.1, 0.1]

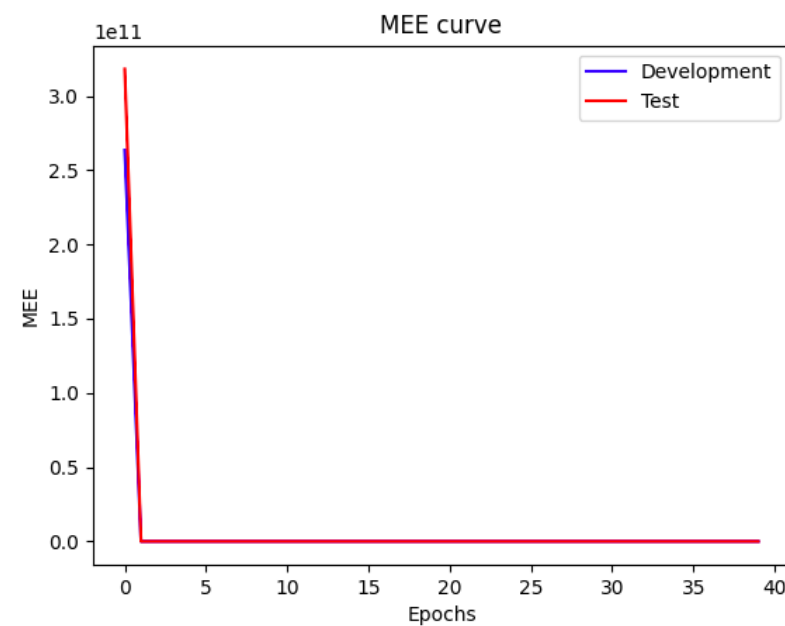
### Model 1




### Model 2




### Model 5




### Model 3




### Model 4



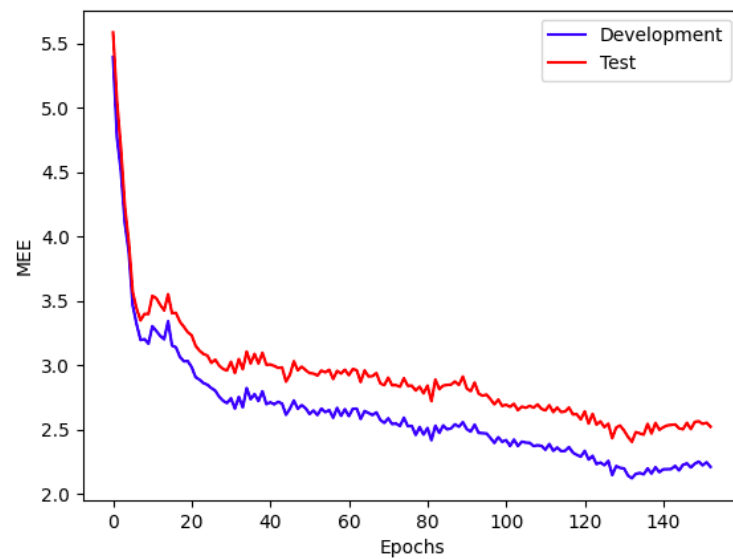
### Model 6




### Model 7



### Model 8



### Model 9



### Model 10

